

PH 354: Computational Physics

Instructors:

Prateek Sharma (prateek@iisc.ac.in)

Manish Jain (mjain@iisc.ac.in)

TA:

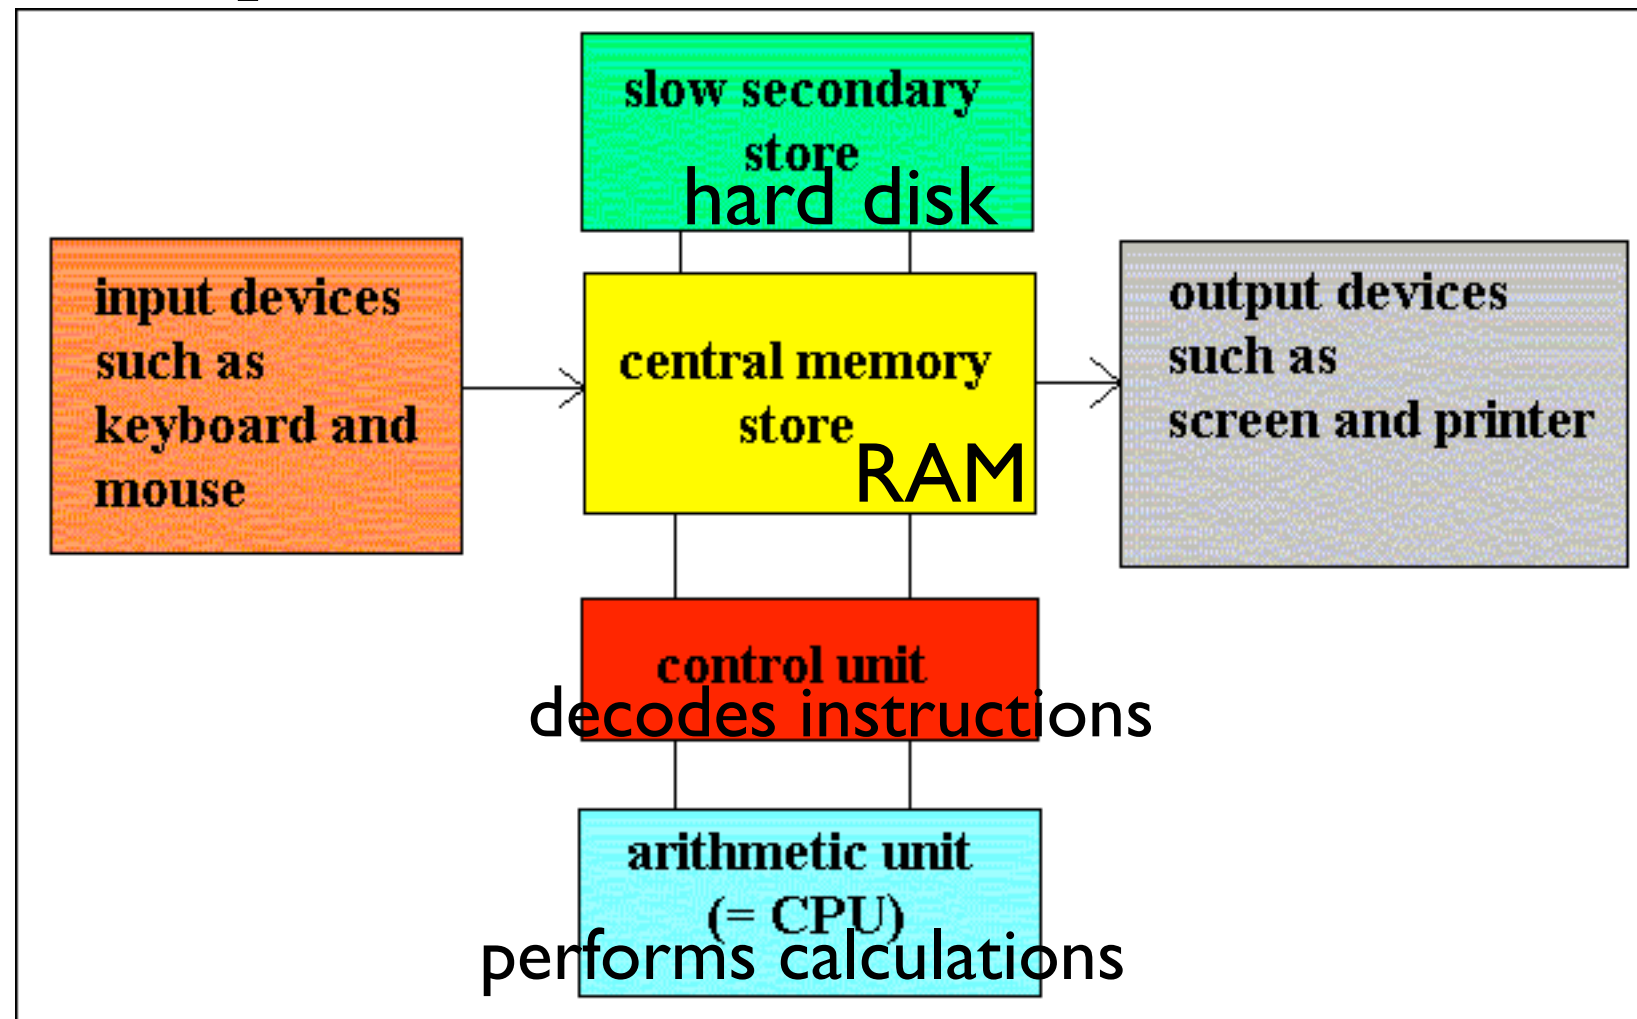
Sanat Gogoi (sanatgogoi@iisc.ac.in)

Philosophy & Approach

- Getting used to write code
- Practise, practise & practise more
- Good programming practices
- Algorithms for Physical applications: root-finding, linear algebra, ODEs, PDEs, etc.
- Interpreted language with little overhead
Python (similar to MATLAB but free & widespread)
- No exams. Heavy HWs & course-project.
Submission via **GitHub**

some basics about how
modern computers
work


Computer Architecture



both instructions & data sent by input devices to memory
loaded from memory to CPU registers

Instruction Set Architecture (ISA): machine language instruction set,
word size, registers

ALU



bitwise logic ops.
AND, OR, NOT, XOR

integer arithmetic ops.
add, subtract, multiply, divide

bit shifting (* or / by 2^n)

FPU: floating point unit

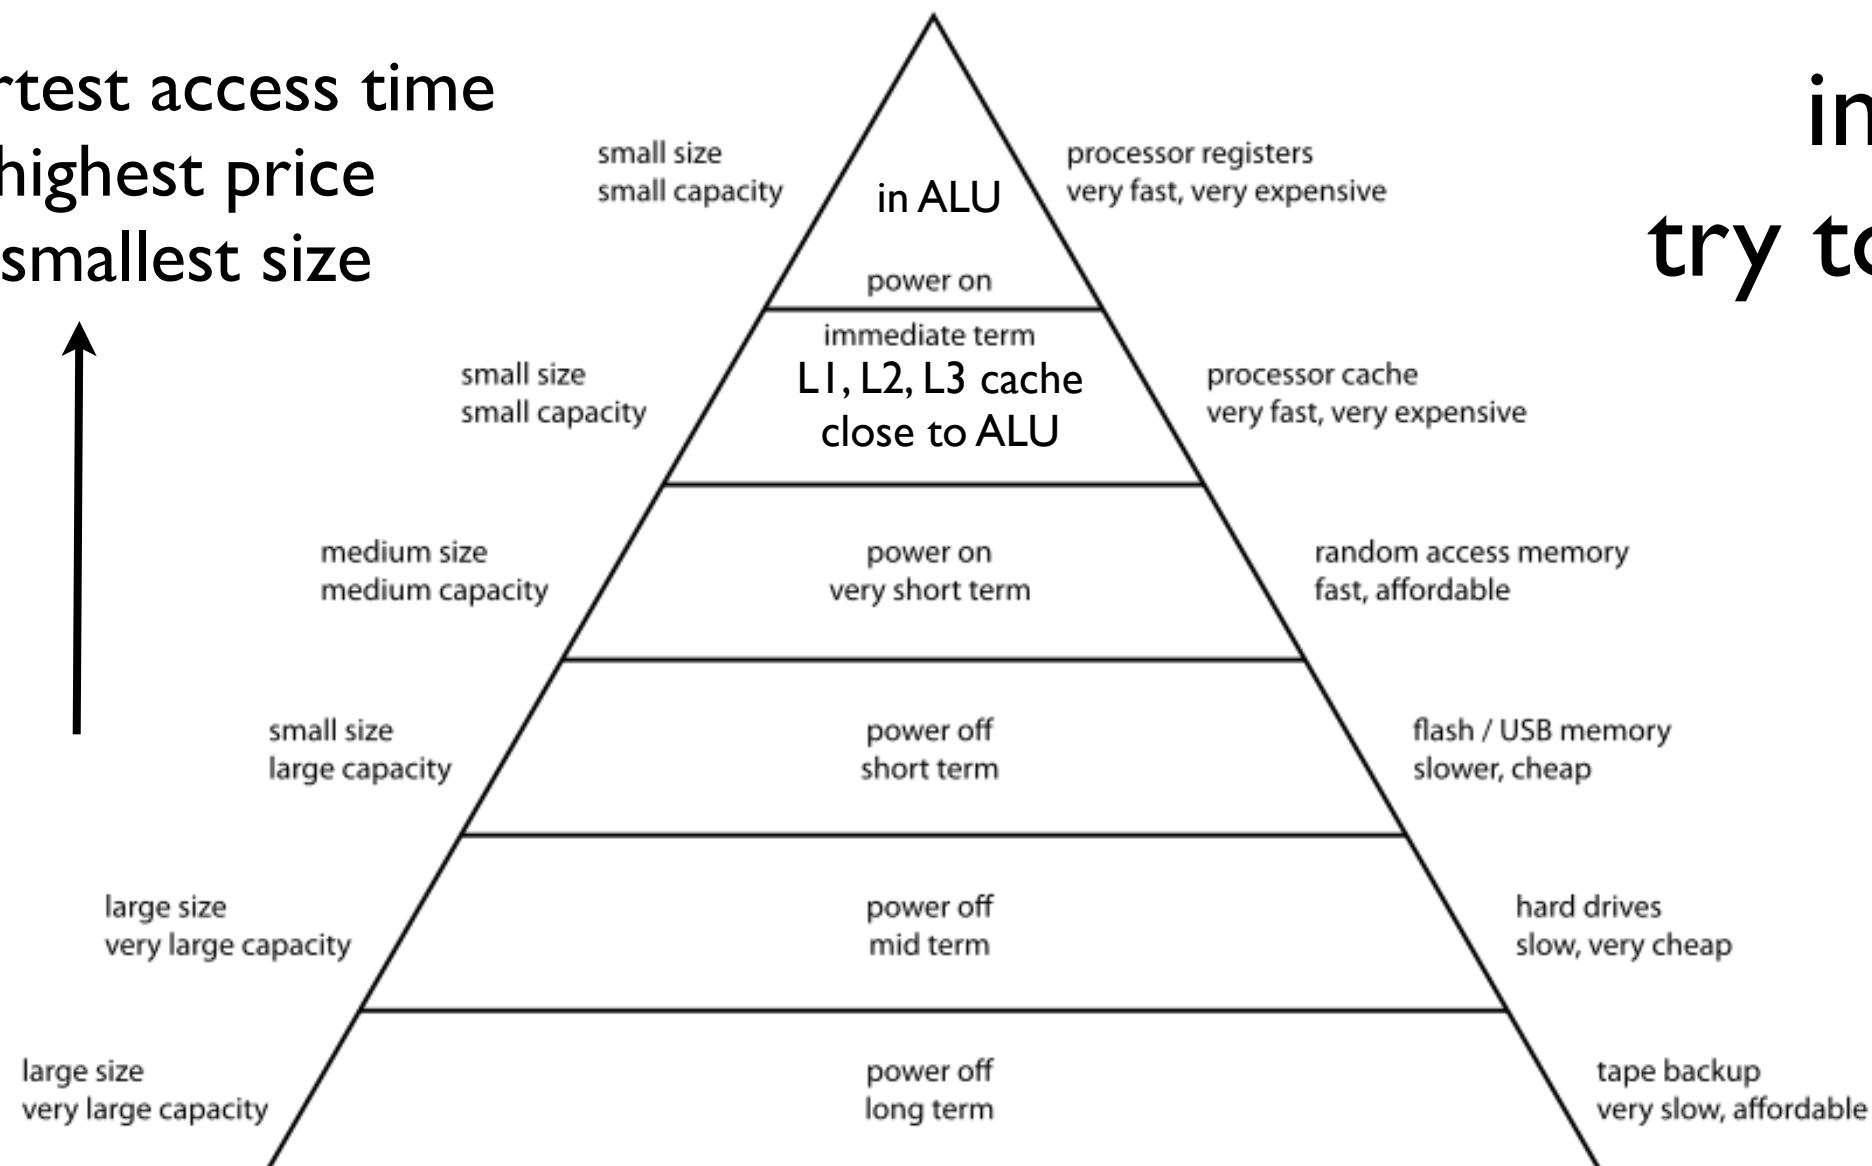
+, -, * fast / slow and so are exp, cos, & other transcendental fns.
commonly used function are coded in machine language

Hierarchical Memory

Computer Memory Hierarchy


shortest access time
highest price
smallest size

implication:
try to reuse cache



Cache Utilization

data stored in memory as a I-D array



column major in fortran

$$\begin{bmatrix} 1 & 2 & 3 \\ 6 & 5 & 4 \end{bmatrix} = [\boxed{1} \boxed{6} \boxed{2} \boxed{5} \boxed{3} \boxed{4}]$$

row major in C!

```
do j = 1, jmax
do i = 1, imax
  a(i,j) = float(i-j)/float(i+j)
enddo
enddo
```



faster than

```
do i = 1, imax
do j = 1, jmax
  a(i,j) = float(i-j)/float(i+j)
enddo
enddo
```

sometimes compilers do these optimizations (-O3)

Latency & Bandwidth


minimum time to do an action
(access time)
rate of action once action is initialized



better to minimize memory access due to latency!


L3 Cache Latency

Latency in ns - Lower is Better




Memory Latency Comparison

ScienceMark 2.0 - Memory Latency in ns (256-byte stride, Lower is Better)



ScienceMark L1 Cache Latency: Time



Nehalem processor:


L1~64 kB

L2~2 MB

L3~30 MB

L1 cache ~ 5 times faster than L3 cache ~5 times faster than RAM!

Clock Rate



clock coordinates
different actions

modern CPUs upto 4 FLOPs per cycle:
 $2.4 \text{ GHz} \Rightarrow 4 \times 2.4 \cdot 10^9 \sim 10^{10} \text{ FLOPs/cycle/core (10 GF)}$
if the cluster has 80 cores \Rightarrow 800 GF machine

this is not the only parameter! since data access is more time-consuming (40 ns) than FLOPs (0.1 ns); having larger RAM/cache/interconnects more important than just clock speed

Architecture level Parallelism

bit level parallelism: 4 bit ... 32 to 64 bit word-size (=register size);
more bits processed/cycle



instruction level parallelism

five-stage pipeline in a RISC (IF = Instruction Fetch, ID = Instruction Decode, EX = Execute, MEM = Memory access, WB = Register write back)


Moore's Law

saturation of clock-speed,
power efficiency, ILP

=> paradigm shift


multiple processors/chip
rather than faster processors

thread/data level parallelism
requires programming
(MPI, openMP) unlike ILP



improvements governed by technology

architecture, compiler, programs reflect this



power issues!

chips becoming smaller and smaller

$$P = \frac{1}{2} C V^2 f$$

higher frequency => more power
consumption & heating
can't be air cooled!

reduce operating voltage (transistor errors),
frequency (speed reduction)

software closely tied to hardware

esp. with parallel systems

source code: high level language (fortran, c, c++)

compiler ↓ (also optimizes the code, e.g., -O2, -O3 flags)

object code & executable (lower level assembly/machine code)

interpreted languages (e.g., python, perl, MATLAB, Mathematica, IDL
scripting languages) slower but handy/easier

important to remember architecture to attain maximum performance

Parallel Computing

multicore: multiple processors on the same chip

shared memory (SMP): all processors have common main memory

Distributed memory: beowulf cluster, parallel clusters w. specialized interconnects


Grid computing: computers communicating over the internet; e.g., SETI@home

GPUs (graphics processing units): driven by games/graphics industry, fast FP operations


Software: MPI (message passing interface; distributed systems),
openMP (shared memory)

distributed memory programming model

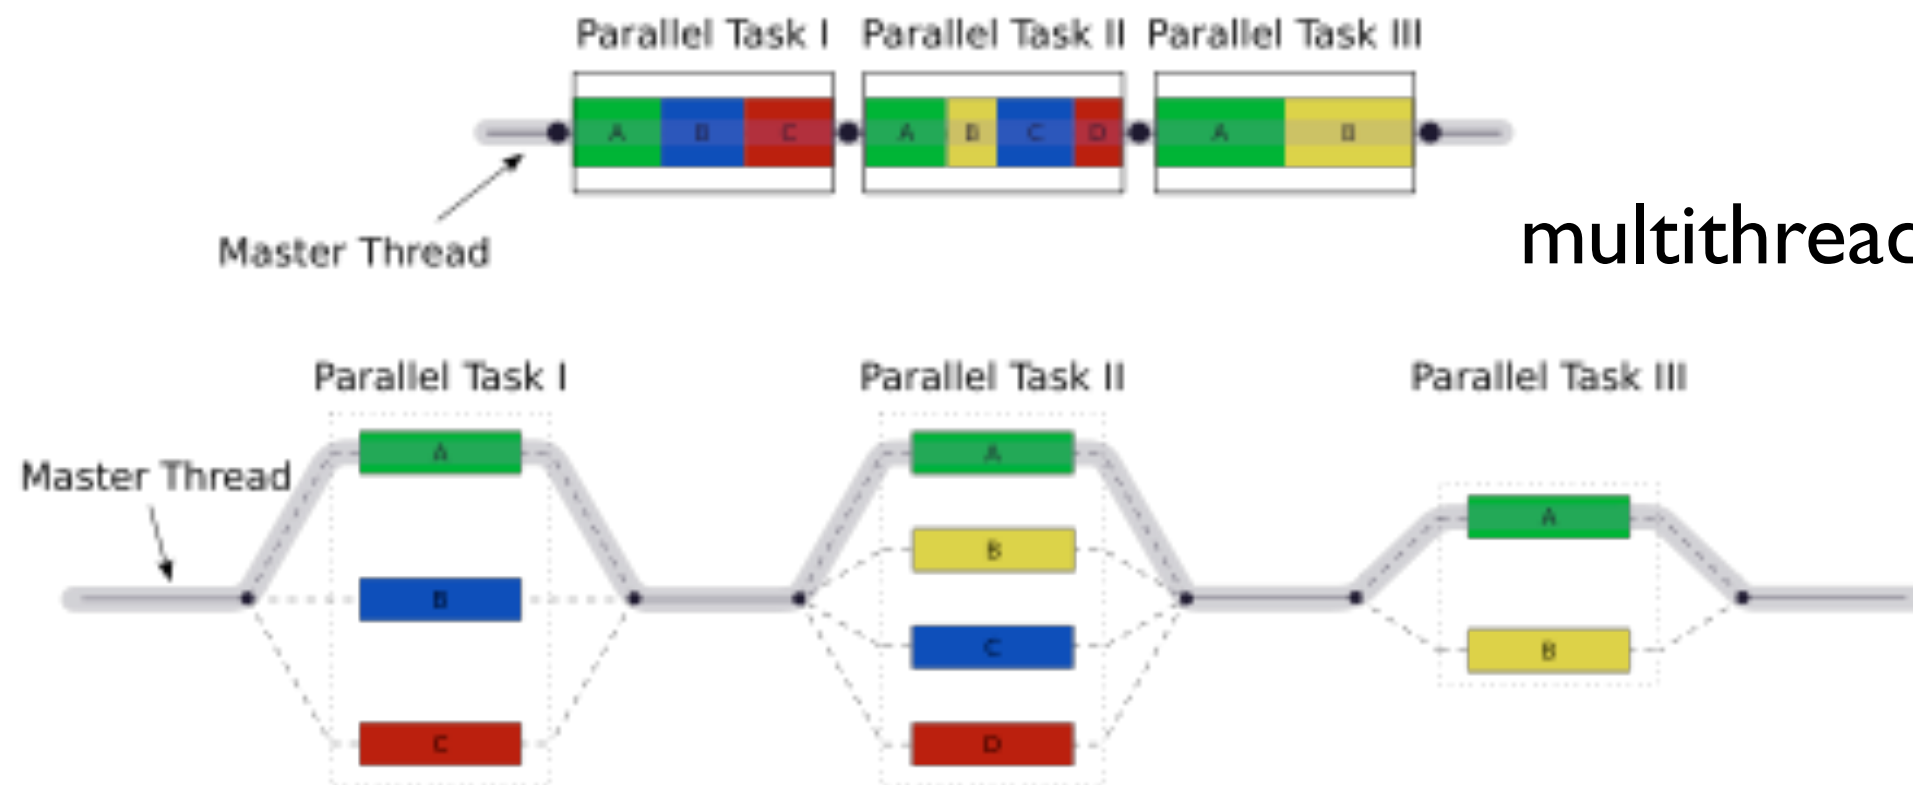
message passing; MPI



Halo Update Communications Pattern



shared memory programming model




multithreading; openMP

Trends in Supercomputing

<http://top500.org/>:

list of 500 fastest (based on LINPACK benchmark) computers in the world
shows trends in architecture, interconnects, vendors, etc.




the list


Rank	Site	Computer
1	RIKEN Advanced Institute for Computational Science (AICS) Japan	K computer, SPARC64 VIIIfx 2.0GHz, Tofu interconnect Fujitsu
2	National Supercomputing Center in Tianjin China	NUDT YH MPP, Xeon X5670 6C 2.93 GHz, NVIDIA 2050 NUDT
3	DOE/SC/Oak Ridge National Laboratory United States	Cray XT5-HE Opteron 6-core 2.6 GHz Cray Inc.
4	National Supercomputing Centre in Shenzhen (NSCS) China	Dawning TC3600 Blade System, Xeon X5650 6C 2.66GHz, Infiniband QDR, NVIDIA 2050 Dawning
5	GSIC Center, Tokyo Institute of Technology Japan	HP ProLiant SL390s G7 Xeon 6C X5670, Nvidia GPU, Linux/Windows NEC/HP
6	DOE/NNSA/LANL/SNL United States	Cray XE6, Opteron 6136 8C 2.40GHz, Custom Cray Inc.
7	NASA/Ames Research Center/NAS United States	SGI Altix ICE 8200EX/8400EX, Xeon HT QC 3.0/Xeon 5570/5670 2.93 Ghz, Infiniband SGI
8	DOE/SC/LBNL/NERSC United States	Cray XE6, Opteron 6172 12C 2.10GHz, Custom Cray Inc.
9	Commissariat a l'Energie Atomique (CEA) France	Bull bullx super-node S6010/S6030 Bull
10	DOE/NNSA/LANL United States	BladeCenter QS22/LS21 Cluster, PowerXCell 8i 3.2 Ghz / Opteron DC 1.8 GHz, Voltaire Infiniband IBM

Some Statistics


Architecture System Share




Interconnect Family System Share




Processor Generation System Share




Operating system Family System Share




**Architecture Share Over Time
1993 - 2011**




**Interconnect Family Share Over Time
1993 - 2011**



**Processor Family Share Over Time
1993 - 2011**



**Countries Share Over Time
1993 - 2011**



testing performance


strong scaling: increase the no. of
procs. on a fixed problem-size

poor scaling for small problem size
communication time \gg computation time
remember latency?

computation $\sim N^3$, communication $\sim N^2$

going to bigger problem size helps with communication overhead

weak scaling: keep the problem per processor the same and inc. the problem-size (& processor count)



Summary

- just touched the tip of the iceberg
- lot of info online
- modern computer architecture:
communication >> computation => cache
contiguous data, minimize data access
- parallel systems: programming models
- free tools (e.g., LAPACK) online; no need
to reinvent the wheel; good to know basics